

PSet6_PF_Evidencias

Está es la estructura del proyecto:

```
project/
  |
  ├── docker-compose.yml
  ├── .env.example
  ├── feature-builder/
  │   ├── Dockerfile
  │   └── build_features.py
  |
  ├── notebooks/
  │   └── 01_ingesta_prices_raw.ipynb
  │       └── 02_build_features_prototipo.ipynb
  │           └── 03_verificacion.ipynb
  |
  └── requirements.txt
```

En jupyter notebook hay un notebook llamado 01_ingesta_prices_raw.ipynb donde se crea el esquema raw, al crearlo en la celda 2, podemos verificar que efectivamente se creó ese esquema metiéndonos en el contenedor y examinando los esquemas, con:

```
docker exec -it postgres bash
psql -U postgres -d marketdb
```

Con lo que ya se ingresa a la base de datos y ahí podemos verlo con

```
\dn
```

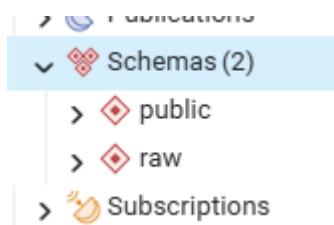
```

PS E:\Octavo Semestre\Data Mining\PSet6> docker exec -it postgres bash
root@36b1d2addd5c:/# psql -U postgres -d marketdb
psql (16.11 (Debian 16.11-1.pgdg13+1))
Type "help" for help.

marketdb=# \dn
      List of schemas
 Schema |        Owner
-----+-----
 public | pg_database_owner
 raw   | postgres

```

En la imagen anterior vemos que se creó el esquema raw, e igualmente podemos verlo en Pgadmin, así:



En el notebook se creó la tabla prices_daily en el esquema raw con las siguientes columnas:

Podemos verificar las columnas y su creación dentro del contenedor de postgres:

```

marketdb=# \dt raw.*
      List of relations
 Schema |        Name        | Type |  Owner
-----+----------------+-----+-----
 raw   | prices_daily | table | postgres
(1 row)

```

Para raw.prices_daily:

```
\d raw.prices_daily
```

Table "raw.prices_daily"					
Column	Type	Collation	Nullable	Default	
date	date		not null		
ticker	character varying(20)		not null		
open	double precision				
high	double precision				
low	double precision				
close	double precision				
adj_close	double precision				
volume	bigint				
run_id	character varying(50)				
ingested_at_utc	timestamp without time zone				
source_name	character varying(50)				
Indexes:					
"prices_daily_pkey" PRIMARY KEY, btree (date, ticker)					

En Pgadmin:



Después de la ingestión de datos en raw.prices_daily, verificamos que se haya insertado el número de filas correcto, que es 5157 filas.

En Pgadmin:

```

1 select count (*) from raw.prices_daily;
2

```

Data Output Messages Notifications

count
bigint

1	5157
---	------

En el contenedor de postgres:

```
root@025a363f06ff:/# psql -U postgres -d marketdb
psql (16.11 (Debian 16.11-1.pgdg13+1))
Type "help" for help.

marketdb=# SELECT COUNT(*) FROM raw.prices_daily;
 count
-----
 5157
(1 row)
```

Creación del esquema analytics y la tabla daily_features

En Pgadmin se corre una sola vez el siguiente código:

```
CREATE SCHEMA IF NOT EXISTS analytics;

CREATE TABLE IF NOT EXISTS analytics.daily_features (
    date DATE NOT NULL,
    ticker VARCHAR(20) NOT NULL,

    -- Identificación temporal
    year INT,
    month INT,
    day_of_week INT,

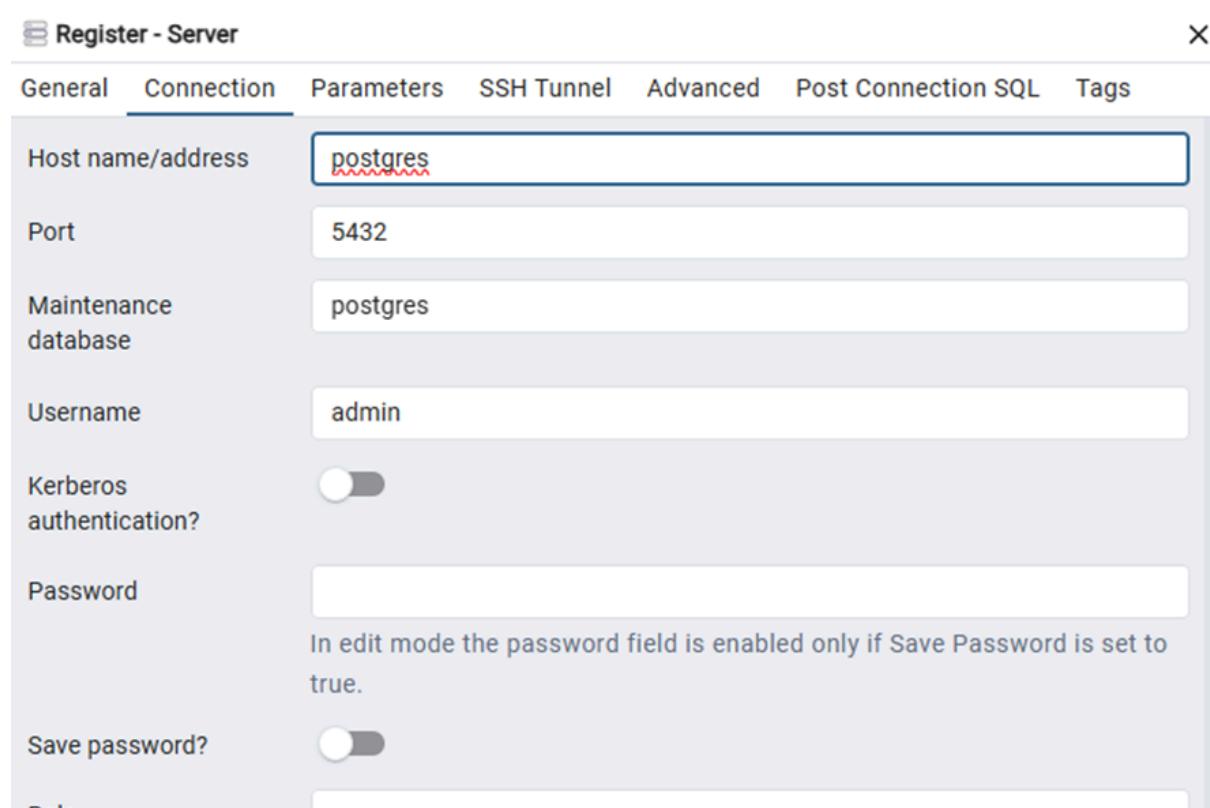
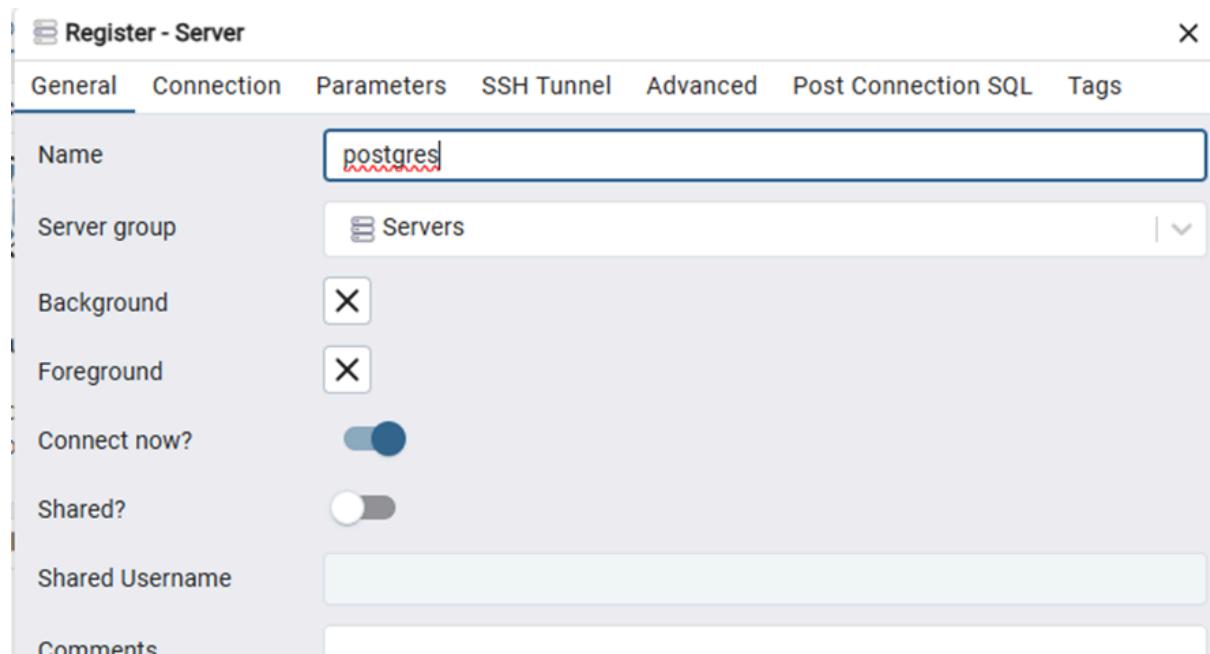
    -- Mercado
    open DOUBLE PRECISION,
    high DOUBLE PRECISION,
    low DOUBLE PRECISION,
    close DOUBLE PRECISION,
    volume BIGINT,

    return_close_open DOUBLE PRECISION,
    return_prev_close DOUBLE PRECISION,
    volatility_5d DOUBLE PRECISION,

    -- Metadatos
    run_id VARCHAR(50),
    ingested_at_utc TIMESTAMPTZ,
```

```
PRIMARY KEY (date, ticker)
);
```

Para la conexión con postgres en Pgadmin:



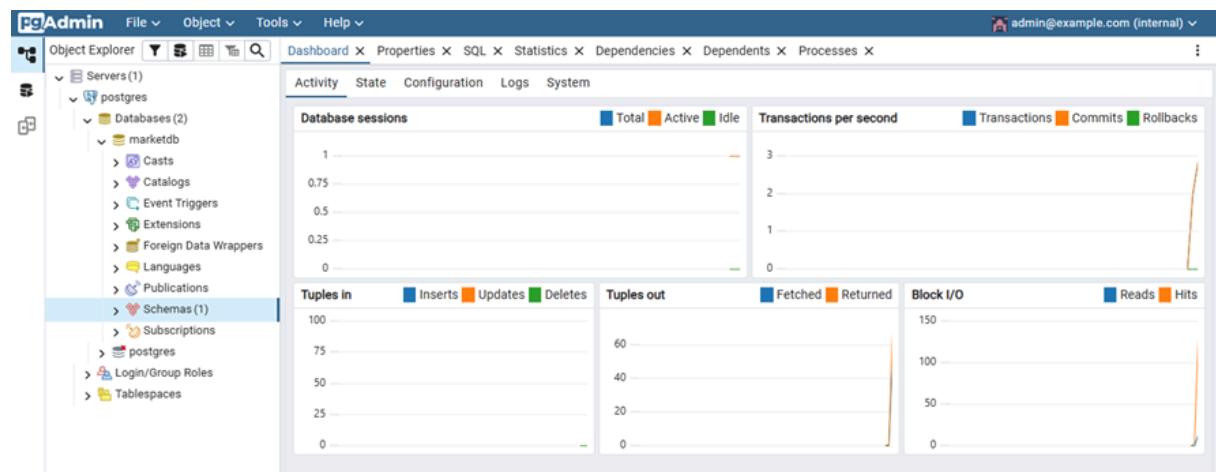
Lo siguiente es lo que tengo en las credenciales del archivo .env

Register - Server

General **Connection** **Parameters** **SSH Tunnel** **Advanced** **Post Connection SQL** **Tags**

Host name/address	postgres
Port	5432
Maintenance database	marketdb
Username	postgres
Kerberos authentication?	<input type="checkbox"/>
Password	*****
In edit mode the password field is enabled only if Save Password is set to true.	
Save password?	<input type="checkbox"/>
Role	

Conexión exitosa de pgadmin con postgres:



Comandos para ingestar a la tabla analytics.daily_features

```
docker compose run feature-builder --mode full --ticker AAPL --overwrite true
docker compose run feature-builder --mode full --ticker MSFT --overwrite true
```

```
docker compose run feature-builder --mode full --ticker TSLA --overwrite t  
rue
```

```
PS E:\Octavo Semestre\Data Mining\PSet6> docker compose run feature-builder --mode full --ticker  
MSFT --overwrite false  
time="2025-11-29T15:40:55-05:00" level=warning msg="Found orphan containers ([pset6-feature-build  
er-run-86e86663a079 pset6-feature-builder-run-252ed8741bdd pset6-feature-builder-run-fbfff605044d7  
pset6-feature-builder-run-a337cfef2501]) for this project. If you removed or renamed this service  
in your compose file, you can run this command with the --remove-orphans flag to clean it up."  
[+] Creating 1/1  
✓ Container postgres Running 0.0s  
Insertando (idempotente)...  
Procesadas 1719 filas para MSFT  
Fecha min: 2019-01-02 00:00:00, max: 2025-10-31 00:00:00
```

Tabla analytics.daily_features

Column	Type	Collation	Nullable	Default
date	timestamp without time zone			
ticker	text			
open	double precision			
high	double precision			
low	double precision			
close	double precision			
adj_close	double precision			
volume	bigint			
run_id	text			
ingested_at_utc	timestamp with time zone			
source_name	text			
year	integer			
month	integer			
day_of_week	integer			
return_close_open	double precision			
return_prev_close	double precision			
volatility_5d	double precision			

En este caso se insertó el ticker Tesla.

The screenshot shows the pgAdmin 4 interface with a connection to 'marketdb/postgres@postgres'. A query is running:

```

1 SELECT *
2 FROM analytics.daily_features
3
4

```

The results table has the following columns:

	date timestamp without time zone	ticker text	open double precision	high double precision	low double precision	close double precision	adj_close double precision	volume bigint
1	2015-01-02 00:00:00	TSLA	14.857999801635742	14.88333206176758	14.21733283996582	14.620667457580566	14.620667457580566	71466000
2	2015-01-05 00:00:00	TSLA	14.303333282470703	14.43333396911621	13.810667037963867	14.005999565124512	14.005999565124512	80527500
3	2015-01-06 00:00:00	TSLA	14.003999710083008	14.279999732971191	13.61400032043457	14.085332870483398	14.085332870483398	93928500
4	2015-01-07 00:00:00	TSLA	14.22333358764648	14.3186674118042	13.985333442687988	14.063332557678223	14.063332557678223	44526000
5	2015-01-08 00:00:00	TSLA	14.187333106994629	14.25333309173584	14.000666618347168	14.041333198547363	14.041333198547363	51637500
6	2015-01-09 00:00:00	TSLA	13.928000450134277	13.998666763305664	13.663999557495117	13.77733325958252	13.77733325958252	70024500
7	2015-01-12 00:00:00	TSLA	13.536666870117188	13.63133351135254	13.283332824707031	13.480667114257812	13.480667114257812	89254500
8	2015-01-13 00:00:00	TSLA	13.554667472839355	13.840666770935059	13.394000053405762	13.616666793823242	13.616666793823242	67159500
9	2015-01-14 00:00:00	TSLA	12.388667106628418	13.01333320617676	12.33333015441895	12.845999717712402	12.845999717712402	173278500
10	2015-01-15 00:00:00	TSLA	12.965999603271484	13.050000190734863	12.666666984558105	12.791333198547363	12.791333198547363	78247500

Total rows: 2725 Query complete 00:00:01.054 CRLF Ln 4, Col 1

Evidencias Funcionamiento de la API

Impresión con sangría

```
{
  "status": "API funcionando correctamente"
}
```

localhost:8000/docs#/default/predict_predict_post

Al colocar todos los parámetros de manera correcta:

GET / Root

POST /predict Predict

Parameters

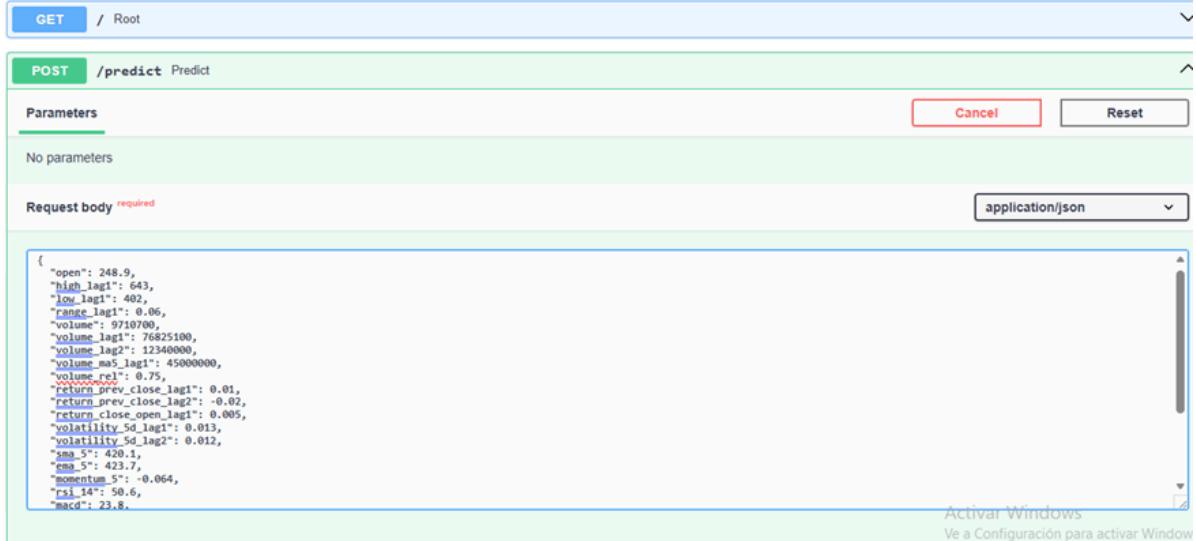
No parameters

Request body required

application/json

```
{ "open": 248.9, "high_lag1": 643, "low_lag1": 40, "range_lag1": 0.06, "volume": 9710700, "volume_lag1": 76825100, "volume_lag2": 12340000, "volume_ma5_lag1": 45000000, "volume_re1": 0.75, "return_prev_close_lag1": 0.01, "return_prev_close_lag2": -0.02, "return_close_open_lag1": 0.005, "volatility_5d_lag1": 0.013, "volatility_5d_lag2": 0.012, "ama_5": 420.1, "ema_5": 421.7, "macd": -0.064, "rsi_14": 50.6, "macd": 23.8 }
```

Activar Windows
Ve a Configuración para activar Window



Curl

```
curl -X 'POST' \
  'http://localhost:8000/predict' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "open": 248.9,
    "high_lag1": 643,
    "low_lag1": 402,
    "range_lag1": 0.06,
    "volume": 9710700,
    "volume_lag1": 76825100,
    "volume_lag2": 12340000,
    "volume_ma5_lag1": 45000000,
    "volume_rel": 0.75,
    "return_prev_close_lag1": 0.01,
    "return_prev_close_lag2": -0.02,
    "return_close_open_lag1": 0.005,
    "volatility_5d_lag1": 0.013,
    "volatility_5d_lag2": 0.012,
    "sma_5": 420.1,
    "ema_5": 423.7,
    "momentum_5": -0.064,
    "rsi_14": 50.6,
    "macd": 23.8,
    "macd_signal": 32.6,
    "boll_position": -0.21,
    "dist_max_5": -0.13,
    "dist_min_5": 0,
    "day_of_week": 3,
    "month": 1
}'
```

Server response

Code	Details
------	---------

200

Response body

```
{  
    "prediction": 1  
}
```

Response headers

```
content-length: 16  
content-type: application/json  
date: Fri, 12 Dec 2025 08:38:11 GMT  
server: uvicorn
```

Responses

Code	Description
------	-------------

200

Successful Response

Si uno de los campos no se incluye:

Curl

```
curl -X 'POST' \
  'http://localhost:8000/predict' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "open": 248.9,
    "high_lag1": 643,
    "low_lag1": 402,
    "range_lag1": 0.06,
    "volume": 9710700,
    "volume_lag1": 76825100,
    "volume_lag2": 12340000,
    "volume_ma5_lag1": 45000000,
    "volume_rel": 0.75,
    "return_prev_close_lag1": 0.01,
    "return_prev_close_lag2": -0.02,
    "return_close_open_lag1": 0.005,
    "volatility_5d_lag1": 0.013,
    "volatility_5d_lag2": 0.012,
    "sma_5": 420.1,
    "momentum_5": -0.064,
    "rsi_14": 50.6,
    "macd": 23.8,
    "macd_signal": 32.6,
    "boll_position": -0.21,
    "dist_max_5": -0.13,
    "dist_min_5": 0,
    "day_of_week": 3,
    "month": 1
  }'
```

Server response